

The Parallelization of Explicit In-time and One Dimensional Finite Difference Solver Using MPI

Mehmet Şamil Dinçer (2236248)

Elvin Gültekinoğlu (2446169)

December 2023

1 Abstract

The Message Passing Interface is an acknowledged standard for distributed computing and it is commonly utilized in different solvers used in several problems. The idea behind MPI is based on dividing the problem into nodes, synchronizing the action of parallel nodes and providing data exchange between different nodes. In that sense, the term Message Passing points out the use of MPI in data transfer. In this paper, the use of MPI in the parallelization of the explicit in-time, one-dimensional finite difference solver is explained in detailed. For the same solver, the implementation of OpenMP is also given with its detail. The performance of MPI implementation is observed by adjusting certain variables such as mesh resolution and number of processors. This analysis is conducted by measuring the duration of the parallelized solver. Another evaluation is performed through strong scaling analysis which is also represented in oncoming sections for both MPI and OpenMP parts. At the end of the report, all results are provided with important points interpreted and according to performance criterion outcomes, the use of MPI and OpenMP are assessed.

2 Introduction

Parallel programming has an important place in engineering applications requiring complex problems which take a considerable amount of time to solve. In that sense, different standards have been set for solving such equations by parallelization. The Message Passing Interface (MPI) is an application program interface which has an extensive use these days. In Bruck et al. (1995) [1], MPI is defined as an industrial standard based on writing portable message-passing parallel programs. In MPI, each parallel unit has its own local memory and data is shared between these units. In one paper, Skjellum et al. (1994) [2] explains MPI as encompassing point-to-point and collective message passing, communication scoping and virtual topologies within the scope of a model of distributed computing. The use of it provides user a flexible, practical and portable interface so that it is convenient to implement MPI in solvers. At this point, it is important to note that it is not a software library rather it is a standard in writing message passing programs and used in related libraries.

The explicit in-time, one dimensional finite difference solver which handles the diffusion equation is equipped with point-to-point communication between processors through MPI. The code implementation is performed on a code file named as "heat1d.c" and after it is compiled, it is run by setting the number of nodes per processor. By doing so, this specified node number is distributed successively to the ranks, which are the indices of current processes among all of them. To evaluate the performance of such application, the number of meshes and number of processors are changed and the computation timings are recorded.

In the next part, the code is changed by implementing OpenMP for parallel processing. This is done on a new file named as "heat1d_omp.c". To evaluate the results, strong scale analysis is performed for different mesh resolutions in both MPI implementation and OpenMP implementation.

Throughout this paper, the implementation of MPI on a solver and its performance parameters are considered and evaluated. In that sense, computational timings for different mesh numbers and number of processors, and results of strong scale analyses are taken as criteria. All these are provided at the end of the paper.

3 The Theory and Methodology

In this section, the theory behind the solver, the methodology followed to implement MPI and the implementation of OpenMP are explained in detail.

4 The Solver

The solver which tackles the diffusion equation given in Equation 1 is an explicit in-time, one-dimensional finite difference solver.

$$\frac{\partial q}{\partial t} = \nabla(k\nabla q) + f(x, t) \quad (1)$$

The first step is to find the coordinates for uniform spacing. This is performed considering the number of nodes per processor distributed sequentially to the ranks which is the input taken from the user. First, calculate uniform grid spacing by Equation 2. Here x_{max} and x_{min} represent domain boundaries, $size$ stands for the number of current processes and n is equal to the input taken from the user.

$$dx = \frac{x_{max} - x_{min}}{size \times n - 1} \quad (2)$$

Then the coordinates for uniform spacing is calculated as follows in Equation 3. In this equation, i represents the index of current coordinate and it is changed from zero to the input taken from the user, which is the number of nodes per processor distributed sequentially to the ranks and $rank$ represents the index of current process. By varying $rank$ values as the code proceeds, coordinates of uniform spacing are calculated. As implied, this is done under a for loop in the code.

$$x[i] = -dx + (i \times dx) + (rank \times dx \times n) \quad (3)$$

It is important to note that, under a for loop coordinates for rank number zero are also computed; however, later in the code, these coordinates for rank zero are not used. This is an important detail since when i and $rank$ are equal to zero, then the coordinate becomes negative, which is contradictory considering x_{min} is zero.

The right hand side of Equation 1 can be approximated by using central difference method at a specific node j as presented in Equation 4.

$$\nabla(k\nabla q) + f(x, t) \approx rhs(q, t) = \frac{k}{dx^2}(q[j-1] - 2q[j] + q[j+1]) + f(x[j], t) \quad (4)$$

The implementation of $rhs(q, t)$ on solver is performed as the Equation 5. Here dt is also calculated as given in Equation 6.

$$q(x, t + dt) = q(x, t) + dt \times rhs(q, t) \quad (5)$$

$$dt = \frac{CFL \times dx \times dx}{k} \quad (6)$$

These are all contributions made additionally to the given plain code in terms of solver. Their implementation on code in C are provided below. All other necessary variables and their calculations for the solver part are already provided on the code.

```

1 double rhs;
2     // UPDATE the solution based on central differantiation.
3     // qn[i] = q[i] + dt*rhs(q,t)
4     // For OpenMP make this loop parallel also
5     for ( int i = 1; i <= n; i++ ){
6         // COMPLETE THIS PART
7         rhs = k * (q[i-1] - 2*q[i] + q[i+1] + source(x[i],time)) / (dx*dx);
8         qn[i] = q[i] + dt*rhs;
9     }

```

4.1 The Implementation of MPI

It is expected to implement MPI on a given code by using point-to-point communication between processors. This part is valid when the size is bigger than one, which amounts to the number of current processes. This is such because when the size is one, end points overlap each other. For point-to-point communication, certain steps are followed. It should be noted that here q represents the field variable and it is the same variable as in Equation 1.

1. When the current process index is equal to the very last process: $q[n]$ is sent to the next rank.
2. When the current process index is not equal to the very first process: $q[0]$ is received from the previous rank.
3. When the current process index is not equal to the very first process again: $q[1]$ is sent to the previous rank.
4. When the current process index is equal to the last process: $q[n+1]$ is received from the next rank.

As seen from Steps 1, 2, 3 and 4, the transfer of data inside the field variable in two ways is realized for the previous and next ranks. By doing so, the integrity between field is provided and solver is applied. In implementing point-to-point communication, to prevent the appearance of certain undesired conditions, non-blocking point-to-point message passing is preferred. These conditions can be classified as follows: There is no available process to post a receive, MPI does not have enough buffer space to store messages and any process cannot be received since no process is sending data. In that sense, `MPI_Isend` and `MPI_Irecv` are used. In addition to these, to make sure that it is safe to use the buffer without any corruption, `MPI_Wait` is also used. The finalized version of the code with MPI is given below.

```

1  # include <math.h>
2  # include <stdlib.h>
3  # include <stdio.h>
4  # include <time.h>
5
6  # define OUT 0
7
8  // Include MPI header
9  # include "mpi.h"
10
11 // Function definitions
12 int main ( int argc, char *argv[] );
13 double boundary_condition ( double x, double time );
14 double initial_condition ( double x, double time );
15 double source ( double x, double time );
16 void runSolver( int n, int rank, int size );
17
18
19
20 /*-----
21    Purpose: Compute number of primes from 1 to N with naive way
22    -----*/
23 // This function is fully implemented for you!!!!!!
24 // usage: mpirun -n 4 heat1d N
25 // N      : Number of nodes per processor
26 int main ( int argc, char *argv[] ){
27     int rank, size;
28     double wtime;
29
30     // Initialize MPI, get size and rank
31     MPI_Init ( &argc, &argv );
32     MPI_Comm_rank ( MPI_COMM_WORLD, &rank );
33     MPI_Comm_size ( MPI_COMM_WORLD, &size );
34
35     // get number of nodes per processor
36     int N = strtol(argv[1], NULL, 10);
37
38
39     // Solve and update the solution in time
40     runSolver(N, rank, size);
41
42     // Terminate MPI.
43     MPI_Finalize ( );
44     // Terminate.
45     return 0;
46 }

```

```

47
48  /*-----
49  Purpose: computes the solution of the heat equation.
50  -----*/
51 void runSolver( int n, int rank, int size ){
52     // CFL Condition is fixed
53     double cfl = 0.5;
54     // Domain boundaries are fixed
55     double x_min=0.0, x_max=1.0;
56     // Diffusion coefficient is fixed
57     double k = 0.002;
58     // Start time and end time are fixed
59     double tstart = 0.0, tend = 10.0;
60
61     // Storage for node coordinates, solution field and next time level values
62     double *x, *q, *qn;
63     // Set the x coordinates of the n nodes padded with +2 ghost nodes.
64     x = ( double*)malloc((n+2)*sizeof(double));
65     q = ( double*)malloc((n+2)*sizeof(double));
66     qn = ( double*)malloc((n+2)*sizeof(double));
67
68     // Write solution field to text file if size==1 only
69     FILE *qfile, *xfile;
70
71     // uniform grid spacing
72     double dx = ( x_max - x_min ) / ( double ) ( size * n - 1 );
73
74     // Set time step size dt <= CFL*h^2/k
75     // and then modify dt to get integer number of steps for tend
76     double dt = cfl*dx*dx/k;
77     int Nsteps = ceil(( tend - tstart )/dt);
78     dt = ( tend - tstart )/(( double )(Nsteps));
79     //printf("nsteps: %d dt: %f \n ",Nsteps, dt);
80     //fflush(stdout);
81
82     int tag;
83     MPI_Status status;
84     double time, time_new, wtime;
85
86     // find the coordinates for uniform spacing
87     for ( int i = 0; i <= n + 1; i++ ){
88         // COMPLETE THIS PART
89         // x[i] = ...
90         x[i] = -dx + ( i*dx ) + (rank * dx * n) ;
91         //printf("rank:%d, i:%d, position:%f \n",rank,i,x[i]); // control position
92
93     }

```

```

94
95 // Set the values of q at the initial time.
96 time = tstart; q[0] = 0.0; q[n+1] = 0.0;
97 for (int i = 1; i <= n; i++){
98     q[i] = initial_condition(x[i],time);
99 }
100
101
102 // Record the starting time.
103 wtime = MPI_Wtime();
104
105
106 // Compute the values of H at the next time, based on current data.
107 for ( int step = 1; step <= Nsteps; step++){
108
109     time_new = time + step*dt;
110
111
112 // Perform point to point communications here!!!!
113
114 if(size>1){
115
116 if (rank != size - 1) {
117     MPI_Request send_right_request;
118     MPI_Isend(&q[n], 1, MPI_DOUBLE, rank + 1, rank, MPI_COMM_WORLD,
119         ↪ &send_right_request);
119     //printf("rank : %d,send to %d, message %f \n",rank,rank+1,q[n]);
120     //fflush(stdout);
121 }
122
123 if (rank != 0) {
124     //double rec;
125     MPI_Request recv_left_request;
126     MPI_Irecv(&q[0], 1, MPI_DOUBLE, rank - 1, rank-1, MPI_COMM_WORLD,
127         ↪ &recv_left_request);
127     MPI_Wait(&recv_left_request, MPI_STATUS_IGNORE); // Wait for the receive
128         ↪ operation to complete
128     //printf("rank : %d,recito %d message %f \n",rank,rank-1,rec);
129     //fflush(stdout);
130     //q[0]=rec;
131 }
132 }
133
134 if (rank != 0) {
135     MPI_Request send_left_request;
136     MPI_Isend(&q[1], 1, MPI_DOUBLE, rank - 1, rank+5, MPI_COMM_WORLD,
137         ↪ &send_left_request);

```

```

137     //printf("rank : %d,send to %d \n",rank,rank+1);
138     //printf("aaa %d",rank);
139 }
140
141 if (rank != size - 1) {
142
143     MPI_Request recv_right_request;
144     MPI_Irecv(&q[n + 1], 1, MPI_DOUBLE, rank + 1, rank+1+5, MPI_COMM_WORLD,
145             ↪ &recv_right_request);
146     MPI_Wait(&recv_right_request, MPI_STATUS_IGNORE); // Wait for the receive
147             ↪ operation to complete
148     //printf("rank : %d,recito %d \n",rank,rank-1);
149 }
150
151 /* // Wait for the completion of the first round of communication
152     MPI_Wait(&send_right_request, MPI_STATUS_IGNORE);
153     MPI_Wait(&recv_left_request, &recv_left_status);
154
155     // Wait for the completion of the second round of communication
156     MPI_Wait(&send_left_request, MPI_STATUS_IGNORE);
157     MPI_Wait(&recv_right_request, &recv_right_status);
158 */
159 }
160
161 double rhs;
162 // UPDATE the solution based on central differantiation.
163 // qn[i] = q[i] + dt*rhs(q,t)
164 // For OpenMP make this loop parallel also
165 for ( int i = 1; i <= n; i++ ){
166     // COMPLETE THIS PART
167     rhs = k * (q[i-1] - 2*q[i] + q[i+1] + source(x[i],time)) / (dx*dx);
168     qn[i] = q[i] + dt*rhs;
169 }
170
171 // q at the extreme left and right boundaries was incorrectly computed
172 // using the differential equation.
173 // Replace that calculation by the boundary conditions.
174 // global left endpoint
175 if (rank==0){
176     qn[1] = boundary_condition ( x[1], time_new );
177 }
178 // global right endpoint
179 if (rank == size - 1 ){
180     qn[n] = boundary_condition ( x[n], time_new );
181 }

```



```

182
183 // Update time and field.
184 time = time_new;
185 // For OpenMP make this loop parallel also
186 for ( int i = 1; i <= n; i++ ){
187     q[i] = qn[i];
188 }
189
190 // In single processor mode, add current solution data to output file.
191 if (size == 1 && OUT==1){
192     for ( int i = 1; i <= n; i++ ){
193         fprintf ( qfile, " %f", q[i] );
194     }
195     fprintf ( qfile, "\n" );
196 }
197
198
199
200
201 if (step==Nsteps){
202     char x_filename[40];
203     char q_filename[40];
204     sprintf(x_filename, "mpi_x_data_size%d_rank%d.txt",size, rank);
205     sprintf(q_filename, "mpi_q_data_size%d_rank%d.txt",size, rank);
206     // write out the x coordinates for display.
207     xfile = fopen ( x_filename, "w" );
208     for (int i = 1; i<(n+1); i++ ){
209         fprintf ( xfile, " %f", x[i] );
210     }
211     fprintf ( xfile, "\n" );
212     fclose ( xfile );
213     // write out the initial solution for display.
214     qfile = fopen ( q_filename, "w" );
215     for ( int i = 1; i < (n+1); i++ ){
216         fprintf ( qfile, " %f", q[i] );
217     }
218     fprintf ( qfile, "\n" );
219     fclose(qfile);
220 }
221 /*
222     if (step==1){
223         char x_filename[40];
224         char q_filename[40];
225         sprintf(x_filename, "x_data_rank%d_step%d.txt", rank,step);
226         sprintf(q_filename, "q_data_rank%d_step%d.txt", rank,step);
227         // write out the x coordinates for display.
228         xfile = fopen ( x_filename, "w" );

```

```

229     for (int i = 1; i<(n+1); i++){
230         fprintf ( xfile, " %f", x[i] );
231     }
232     fprintf ( xfile, "\n" );
233     fclose ( xfile );
234     // write out the initial solution for display.
235     qfile = fopen ( q_filename, "w" );
236     for ( int i = 1; i < (n+1); i++ ){
237         fprintf ( qfile, " %f", q[i] );
238     }
239     fprintf ( qfile, "\n" );
240     fclose(qfile);
241 }
242 */
243
244
245 }
246
247 /*
248     char x_filename[40];
249     char q_filename[40];
250     sprintf(x_filename, "mpi_x_data_rank%d_step%d.txt", rank,10);
251     sprintf(q_filename, "mpi_q_data_rank%d_step%d.txt", rank,10);
252     // write out the x coordinates for display.
253     xfile = fopen ( x_filename, "w" );
254     for (int i = 1; i<(n+1); i++){
255         fprintf ( xfile, " %f", x[i] );
256     }
257     fprintf ( xfile, "\n" );
258     fclose ( xfile );
259     // write out the initial solution for display.
260     qfile = fopen ( q_filename, "w" );
261     for ( int i = 1; i < (n+1); i++ ){
262         fprintf ( qfile, " %f", q[i] );
263     }
264     fprintf ( qfile, "\n" );
265     fclose(qfile);
266 */
267
268 // Record the final time.
269 // if (rank == 0 ){
270 wtime = MPI_Wtime() - wtime;
271
272 // Add local number of primes
273 double global_time = 0.0;
274 MPI_Reduce( &wtime, &global_time, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
275

```

```

276 if(rank==0)
277     printf ( " SIZE: %d RANK: %d Wall clock elapsed seconds = %f\n",size,
    ↪ rank,global_time );
278
279
280 if( size == 1 && OUT==1)
281     fclose ( qfile );
282
283 free(q); free(qn); free(x);
284
285 return;
286 }
287 /*-----*/
288 double boundary_condition ( double x, double time ){
289     double value;
290
291     // Left condition:
292     if ( x < 0.5 ){
293         value = 100.0 + 10.0 * sin ( time );
294     }else{
295         value = 75.0;
296     }
297     return value;
298 }
299 /*-----*/
300 double initial_condition ( double x, double time ){
301     double value;
302     value = 95.0;
303
304     return value;
305 }
306 /*-----*/
307 double source ( double x, double time ){
308     double value;
309
310     value = 0.0;
311
312     return value;
313 }

```

The name of the file containing this source code is "heat1d.c". To compile and run the code, the following lines should be typed on terminal.

- To compile the code: `mpicc -g -Wall -o heat1d heat1d.c -lm`
- To run the code: `mpiexec -n 4 ./heat1d N`

After running the code, the same number of files are created as the number of cores. To check the accuracy of results, a Python script which plots the solution is created and the trends in plots are evaluated.




Figure 1: MPI OUTPUT

4.2 The Implementation of OpenMP

In order to parallelize the code by using OpenMP, the loop inside the code which is responsible for the calculation of $\text{rhs}(q,t)$ is modified. First the required number of steps is calculated by applying the Equation 7, where dt represents the time step and calculated by using CFL value. This is used in the code file with a built-in function, which finds the nearest integer value of a certain number.

$$NSteps = \frac{t_{end} - t_{start}}{dt} \quad (7)$$

For all these steps, a loop is created to perform calculation one by one with two other interior loops. The purpose of these interior loops is to initiate the related solver part, which is $\text{rhs}(q,t)$. An important point to be careful is the memory since while executing a parallel region, OpenMP threads share the same memory. In order to prevent undesired results because of this condition, shared memory model is utilized. Variables outside these for loops are specified as shared, and other loop variables are indicated as private. The part in which OpenMP is used in the code is given as the following.

```
1  for ( int step = 1; step <= Nsteps; step++ ){
2      time_new = time + step*dt;
3      double rhs;
4      // UPDATE the solution based on central differantiation.
5      // qn[i] = q[i] + dt*rhs(q,t)
6      // For OpenMP make this loop parallel also
7      int i;
8      #pragma omp parallel default(none) private(i, rhs) shared(size, n, q, k, x, dx,
9      ↪ qn, dt, time)
10     {
11         #pragma omp for
12         for (i = 1; i <= (size * n) - 2; i++) {
13             // COMPLETE THIS PART
14             rhs = k * (q[i - 1] - 2 * q[i] + q[i + 1] + source(x[i], time)) / (dx *
15             ↪ dx);
16             qn[i] = q[i] + dt * rhs;
17         }
18     }
19
20     // q at the extreme left and right boundaries was incorrectly computed
21     // using the differential equation.
22     // Replace that calculation by the boundary conditions.
23     // global left endpoint
24     qn[0] = boundary_condition ( x[0], time_new );
25     // global right endpoint
26     qn[size*n-1] = boundary_condition ( x[size*n-1], time_new );
27     // Update time and field.
```

```

27     time = time_new;
28     // For OpenMP make this loop parallel also
29     #pragma omp parallel default(none) private(i) shared(n, size, q, qn)
30     {
31         #pragma omp for
32         for (i = 0; i <= n * size - 1; i++) {
33             q[i] = qn[i];
34         }
35     }
36 }

```

The name of the file containing this source code is "heat1d_omp.c". To compile and run the code, the following lines should be typed on terminal.

- To compile the code: `gcc -fopenmp -o heat1d_omp heat1d_omp.c -lm`
- To run the code: `./heat1d_omp N 4`

After running the code, the same number of files are created as the number of cores. To check the accuracy of results, a Python script which plots the solution is created and the trends in plots are evaluated.

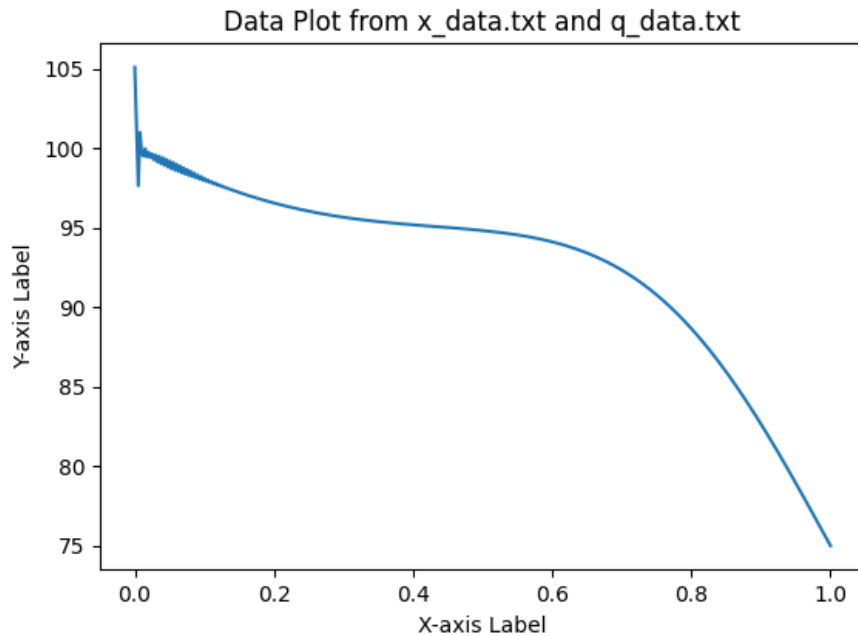


Figure 2: OpenMP OUTPUT

5 Results

In this section, to measure the performance of both MPI and OpenMP, several findings are provided.

5.1 Results of Mesh Resolution

For different mesh resolutions, both version of codes, MPI and OpenMP are run and timings are recorded by keeping the core number as the same. As observed from Figure 3, both implementations have the uprising trend as the mesh resolution increases. This is expected since more mesh resolution takes more time to solve. Another observation is that using MPI takes less time than using OpenMP so it can be said that OpenMP is slower.

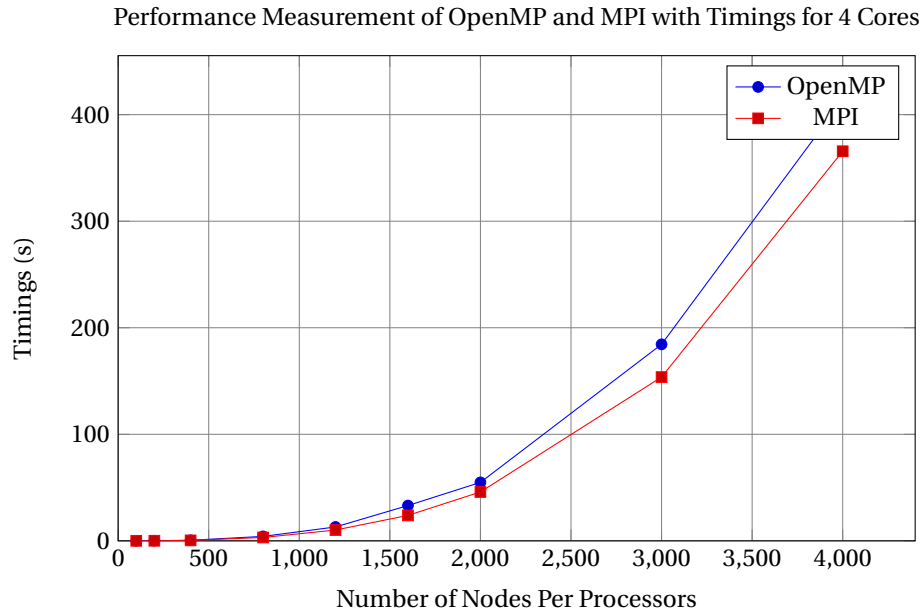


Figure 3: Performance Measurement of OpenMP and MPI with Timings for Fix Core Number(4)

5.2 Results of Comparison the Number of Processes

Similar to the trend in 3, for varying number of processes, both MPI and OpenMP have the similar trend, which is uprising. However, as the number of processes increase, the gap between these two also increases in terms of timings. Considering this phenomena, OpenMP is slower here also.

Performance Measurement of OpenMP and MPI with Timings for Fix Node Number Per Processor(1000)




Figure 4: Performance Measurement of OpenMP and MPI with Timings for the Changing Processor Number

5.3 Strong Scaling Analysis

In strong scaling study, the problem size is kept fixed and the time for different number of processes is measured. As expected, increasing the number of processes decreases the required time since the node numbers per processor decreases. Here it is also observed that OpenMP is slower than MPI.

Performance Measurement of OpenMP and MPI with Timings for Fix Problem(12000 Nodes Calculation)




Figure 5: Performance Measurement of OpenMP and MPI with Timings for Strong Scaling Analysis

6 Comments and Conclusion

As seen in figures 4, the timing increases dramatically with the number of nodes per cores. This is the expected result. Because when Nstep is calculated, sqrt of (the total core number*node per core) is used.

As seen in figures 4, the timing increases dramatically with the number of cores. This is the expected result. Same thing happen in this calculation too. Because when Nstep is calculated, sqrt of (the total core number*node per core) is used.

As seen in figures 5, the timing decreases with number of core for fixed problem. But this is not fully linear because of unparallel part and communication.

References

- [1] Jehoshua Bruck et al. “Efficient message passing interface (MPI) for parallel computing on clusters of workstations”. In: *Proceedings of the seventh annual ACM symposium on Parallel algorithms and architectures*. 1995, pp. 64–73.
- [2] Anthony Skjellum et al. “Extending the message passing interface (MPI)”. In: *Proceedings Scalable Parallel Libraries Conference*. IEEE. 1994, pp. 106–118.